

Zheyuan Ryan Shi

Last updated: May 2021

CONTACT	Ryan Shi (preferred name) ISR - TCS Hall 430, 4665 Forbes Avenue Pittsburgh, PA 15213, USA	ryanshi@cmu.edu https://ryanzshi.github.io
EDUCATION	School of Computer Science, Carnegie Mellon University, USA Ph.D. Candidate, Societal Computing.	Aug 2018 - Present
	Swarthmore College, USA B.A. Mathematics and Computer Science with Honors, 4.00/4.00	Aug 2014 - May 2018
	Massachusetts Institute of Technology, USA Visiting Student, EECS, 5.00/5.00	Aug 2016 - May 2017
EMPLOYMENT	Machine Learning Engineer Intern, Facebook Feed Integrity.	May 2021 - Aug 2021
	Summer Analyst, Credit Suisse Credit and equity derivative structuring, quantitative modeling group.	Jun 2017 - Aug 2017
RESEARCH AREAS	Machine learning, online learning, reinforcement learning, game theory. AI for social good, as related to social work, security and sustainability. Deployment and evaluation of AI in non-profit settings.	
TECHNICAL SKILLS	Python (PyTorch, Scikit-learn, Pandas, Matplotlib, etc.), CPLEX, Gurobi, SQL, Matlab.	
DEPLOYED RESEARCH	Data-driven Optimization of Volunteer-based Community Services Mobile app notification scheme adopted and deployed by 412 Food Rescue, since January 2020. Deep learning-based recommender system improved hit rate from 45% to 79%.	
	CyberTWEAK: Personalized Solution for Watering Hole Attacks Chrome extension, hosted on AWS, publicly available at bit.ly/CyberTWEAK .	
WORKING PAPERS	Zheyuan Ryan Shi , Claire Wang, Fei Fang. Artificial Intelligence for Social Good: A Survey. Under review, 2020. Available on arXiv.	
	Zheyuan Ryan Shi , Zhiwei Steven Wu, Rayid Ghani, Fei Fang. Bandit Data-driven Optimization: AI for Social Good and Beyond. Work in progress, 2020. Available on arXiv.	
RIGOROUSLY REFEREED CONFERENCE PUBLICATIONS	Zheyuan Ryan Shi , Leah Lizarondo, Fei Fang. A Recommender System for Crowdsourcing Food Rescue Platforms. In Proceedings of the Thirtieth The Web Conference (WWW-21).	
	Zheyuan Ryan Shi . AI for Social Good: Between My Research and the Real World. In Proceedings of the Thirty-Fifth AAAI Conference on Artificial Intelligence, Doctoral Consortium (AAAI-21-DC).	
	Zheyuan Ryan Shi , Ariel D. Procaccia, Kevin S. Chan, Sridhar Venkatesan, Noam Ben-Asher, Nandi O. Leslie, Charles Kamhoua, Fei Fang. Learning and Planning in the Feature Deception Problem. In Proceedings of the Eleventh Conference on Decision and Game Theory for Security (GameSec-20).	

Zheyuan Ryan Shi*, Yiwen Yuan*, Kimberly Lo, Leah Lizarondo, Fei Fang. Improving Efficiency of Volunteer-Based Food Rescue Operations. In Proceedings of the Thirty-Second Annual Conference on Innovative Applications of Artificial Intelligence (IAAI-20).

Zheyuan Ryan Shi, Aaron Schlenker, Brian Hay, Daniel Bittleston, Siyu Gao, Emily Peterson, John Trezza, Fei Fang. Draining the Water-hole: Mitigating Social Engineering Attacks with CyberTWEAK. In Proceedings of the Thirty-Second Annual Conference on Innovative Applications of Artificial Intelligence (IAAI-20).

Yufei Wang, **Zheyuan Ryan Shi**, Lantao Yu, Yi Wu, Rohit Singh, Lucas Joppa, Fei Fang. Deep Reinforcement Learning for Green Security Games with Real-Time Information. In Proceedings of the Thirty-Third AAAI Conference on Artificial Intelligence (AAAI-19). Selected for oral.

Zheyuan Ryan Shi*, Ziye Tang*, Long Tran-Thanh, Rohit Singh, Fei Fang. Designing the Game to Play: Optimizing Payoff Structure in Security Games. In Proceedings of the 27th International Joint Conference on Artificial Intelligence and the 23rd European Conference on Artificial Intelligence (IJCAI-ECAI-18).

REFEREED
WORKSHOP
PUBLICATIONS

Zheyuan Ryan Shi, Zhiwei Steven Wu, Rayid Ghani, Fei Fang. Bandit Data-driven Optimization: AI for Social Good and Beyond. At the Machine Learning for Economic Policy Workshop at NeurIPS-20.

Zheyuan Ryan Shi, Zhiwei Steven Wu, Rayid Ghani, Fei Fang. Bandit Data-driven Optimization: AI for Social Good and Beyond. At the Consequential Decisions in Dynamic Environments Workshop at NeurIPS-20.

Zheyuan Ryan Shi, Zhiwei Steven Wu, Rayid Ghani, Fei Fang. Bandit Data-driven Optimization: AI for Social Good and Beyond. At the Machine Learning for the Developing World Workshop at NeurIPS-20.

Zihan Zhou, **Zheyuan Ryan Shi**, Fei Fang, Yi Wu. Approximated Temporal-Induced Neural Self-Play for Finitely Repeated Bayesian Games. At the Workshop on Reinforcement Learning in Games at AAAI-20.

Yiwen Yuan, Kimberly Lo, **Zheyuan Ryan Shi**, Leah Lizarondo, Fei Fang. Efficiency and Fairness of Food Rescue Platforms: An Initial Study. At the AI for Social Good Workshop at IJCAI-19.

Zheyuan Ryan Shi, Ariel D. Procaccia, Kevin S. Chan, Sridhar Venkatesan, Noam Ben-Asher, Nandi O. Leslie, Charles Kamhoua, Fei Fang. Feature Deception Games. At the Strategic Reasoning Workshop at IJCAI-19.

Zheyuan Ryan Shi, Ariel D. Procaccia, Kevin S. Chan, Sridhar Venkatesan, Noam Ben-Asher, Nandi O. Leslie, Charles Kamhoua, Fei Fang. Learning and Planning in Feature Deception Games. At the Machine Learning in the Presence of Strategic Behavior Workshop at EC-19.

Zheyuan Ryan Shi, Aaron Schlenker, Brian Hay, Fei Fang. Draining the Water-hole: Mitigating Social Engineering Attacks. At the Artificial Intelligence for Cyber Security (AICS) Workshop at AAAI-19.

Yufei Wang, **Zheyuan Ryan Shi**, Lantao Yu, Yi Wu, Rohit Singh, Lucas Joppa, Fei Fang. Deep Reinforcement Learning for Green Security Games with Real-Time Information. At the Reinforcement Learning in Games Workshop at AAAI-19.

Lantao Yu, Yi Wu, **Zheyuan Ryan Shi**, Rohit Singh, Lucas Joppa, Fei Fang. Deep Reinforcement

Learning for Green Security Games with Real-Time Information. At the AI for Wildlife Conservation (AIWC) Workshop at IJCAI-18.

Zheyuan Ryan Shi*, Ziyi Tang*, Long Tran-Thanh, Rohit Singh, Fei Fang. Designing the Game to Play: Optimizing Payoff Structure in Security Games. At the AI for Wildlife Conservation (AIWC) Workshop at IJCAI-18.

Zheyuan Ryan Shi*, Ziyi Tang*, Long Tran-Thanh, Rohit Singh, Fei Fang. Designing the Game to Play: Optimizing Payoff Structure in Security Games. At the International Workshop on Optimization in Multiagent Systems (OptMAS-18) at AAMAS-18.

Dhaval Adjodah, Yan Leng, Shi Kai Chong, **Zheyuan Ryan Shi**, Peter Krafft, Alejandro Noriega, Sandy Pentland. Social Bayesian Decision Making. At the 52nd Conference on Information Sciences and Systems (CISS-18).

Zheyuan Ryan Shi, Fei Fang. Optimizing Peer Teaching to Enhance Team Performance. In Autonomous Agents and Multiagent Systems: AAMAS'17 Workshops Best Papers, Volume 10642 of Lecture Notes in Artificial Intelligence, Springer, 2017. Winner of Best Paper at TEAMAS-17.

OTHER
REFEREED
PUBLICATIONS

Zheyuan Ryan Shi, Sindhu Kutty. Strategic Reporting in Exponential Family Prediction Markets. In Proceedings of the 2016 MIT IEEE Undergraduate Research Technology Conference (IEEE URTC 2016).

HONORS AND
AWARDS

Upsilon Pi Epsilon Honor Society Scholarship, IEEE Computer Society, 2018
Nomination to Phi Beta Kappa, 2018
Best Paper Award, First International Workshop on Teams in Multiagent Systems (TEAMAS), 2017
Honorable Mention, CRA Outstanding Undergraduate Researcher Award, 2017
Conference travel grants: AAAI-20, IJCAI-ECAI-18, RecSys-16

SERVICES

Conference Organizer

2021 Symposium on AI and Social Good, April 2021.
2020 CMU Symposium on AI and Social Good, April 2020.

Program Committee

ACM COMPASS-21, June 2021.
AAAI-21 Special Track on AI for Social Impact, February 2021.
IJCAI-20 Workshop on AI for Social Good, January 2021.
Harvard CRCS Workshop on AI for Social Good, July 2020.
AAAI-20 Special Track on AI for Social Impact, February 2020.
AAAI-20 Outreach event Try AI, February 2020.
NeurIPS-19 Workshop on AI for Social Good, December 2019.
IJCAI-19 Workshop on AI for Social Good, August 2019.

TEACHING

Teaching Assistant

Artificial intelligence methods for social good (17-737), Carnegie Mellon University, Spring 2021.
Artificial intelligence methods for social good (17-737), Carnegie Mellon University, Spring 2020.
Introduction to Econometrics (ECON 031), Swarthmore College, Spring 2016.

Other teaching positions

Math Clinician for all undergraduate math courses, Swarthmore College, Spring 2016.

RESEARCH
EXPERIENCE

School of Computer Science, Carnegie Mellon University

Sept 2018 - Present

Working towards AI for immediate and measurable social good. Learning in games in cybersecurity

and sustainability domains. Formulated and analyzed two-layer optimization in Stackelberg security games. Applied deep reinforcement learning to solve security games with online information.

Supervisor: Fei Fang

Department of Mathematics, Swarthmore College

Sept 2017 - May 2018

Proposed homological symmetry and analyzed its use in topological data analysis. Experimented with planar curves and handwritten digits.

Supervisor: Noah Giansiracusa

CRCS, Harvard University

Sept 2016 - May 2017

Formulated the peer teaching problem and devised algorithms to boost team performance. Published and chosen as Best Paper in TEAMAS'17. Designed and experimented with doodle polling mechanisms.

Supervisor: Fei Fang

Human Dynamics Group, Media Lab, MIT

Sept 2016 - May 2017

Studied network structure in evolution strategies for deep reinforcement learning. Investigated influence patterns in SciCast prediction markets. Built DRL testing module on Amazon EC2, and Python package for the Influence Model.

Supervisor: Dhaval Adjodah

Department of Computer Science, Swarthmore College

Jan 2016 - Aug 2016

Investigated incentive compatibility in exponential family prediction markets. Published at IEEE URTC'16. Designed and analyzed cryptogenography protocols.

Supervisors: Sindhu Kutty, Joshua Brody

OTHER
EXPERIENCE

The Cornell, Maryland, Max Planck Pre-doctoral Research School

Aug 2018

Selected as one of the 80 attendees worldwide with scholarship.